



WBL Deep Learning:: Part 2

Beate Sick, Oliver Dürr


Convolutional Neural Networks

Topics of today



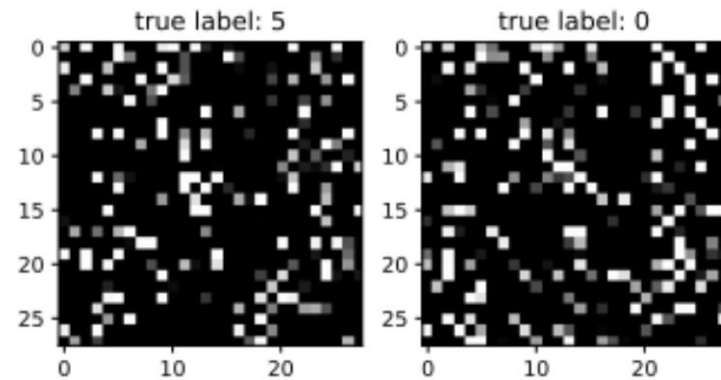



- Convolutional Neural Networks (CNN) for images
 - Motivation of CNN
 - Use local correlation structure of images
 - How does convolution work?
 - Kernel and filter with shared weights
 - Each kernel yields one activation map
 - How does pooling work?
 - Architecture of CNNs



Recall: Architecture of a fully connected NN




MNIST exercise: Does shuffling disturb a fcNN?



→ The performance of a fcNN is the same on original and shuffled images

Recall: Imagenet challenge

1000 classes
1 Mio samples



Human: 5% misclassification


Only one non-CNN approach in 2013

GoogLeNet 6.7%

A. Krizhevsky
first CNN in 2012


Und es hat zoom gemacht

An artificial neuron




Convolution extracts local information using few weights

Fully connected neural net



Locally connected neural net



Shared weights:


by using the **same weights for each patch** of the image we need much **less parameters** than in the fully connected NN and get from each patch the same kind of **local feature information** such as the presence of an edge.

Example of designed Kernel / Filter



$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix}$$

Horizontal Sobel kernel



Applying a vertical edge detector kernel

Convolution

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input X

1	0	1
0	1	0
1	0	1

Kernel W

4	3	4
2	4	3
2	3	4

Result Z


Convolution (let's ignore bias b):

$$z = b + \sum_i x_i w_i$$


1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

4		

CNN Ingredient I: Convolution



Zero-padding to achieve
same size of feature and input



no padding to only use
valid input information

The *same* weights are used at each position of the input image.


Exercise: Do one convolution step by hand



The kernel is 3x3 and is applied at each valid position
– how large is the resulting activation map?

The small numbers in the shaded region are the kernel weights.
Determine the position and the value within the resulting activation map.

3	3	2	1	0
0_0	0_1	1_2	3	1
3_2	1_2	2_0	2	3
2_0	0_1	0_2	2	2
2	0	0	0	1





Convolutional networks use neighborhood information and replicated local feature extraction

In a locally connected network the calculation rule

$$z = b + \sum_i x_i w_i$$



Pixel values in a small image patch are element-wise multiplied with weights of a small filter/kernel:




The filter is applied at each position of the image and it can be shown that the **result is maximal if the image pattern corresponds to the weight pattern.**

The results form again an image called **feature map** (=activation map) which shows at which position the feature is present.

feature/activation map




Convolutional networks use neighborhood information and replicated local feature extraction



The weights of each filter are randomly initiated and then adapted during the training.


Convolution layer with a 1-channel input and 6 kernels



Convolution of the input image with 6 different kernels results in 6 activation maps.
If the input image has only one channel, then each kernel has also only one channel.

Animated convolution with 3 input channels


3 color channel input image



$$z = b + \sum_i x_i w_i$$


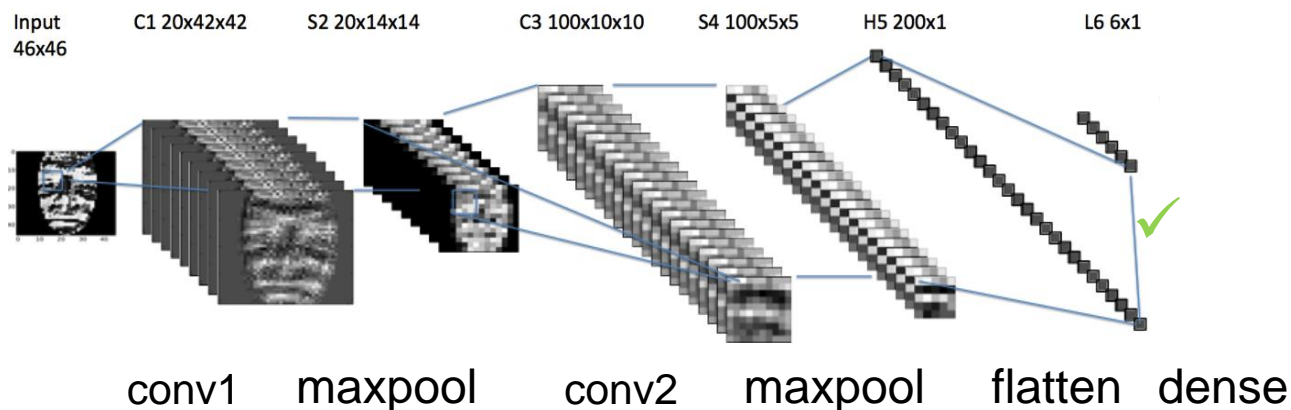
Animation credits: M.Gorner, <https://codelabs.developers.google.com/codelabs/cloud-tensorflow-mnist/#10>
For an example with number see convolution demo in: <https://cs231n.github.io/convolutional-networks/>

Convolution layer with a 3-channel input and 6 kernels



Convolution of the input image with 6 different kernels results in 6 activation maps. If the input image has 3 channels, then each filter has also 3 channels.


CNN ingredient II: Maxpooling Building Blocks reduce size



Simply join e.g. 2x2 adjacent pixels in one by taking the max.
→ less weights in model
→ Less train data needed
→ increased performance


Hinton: „The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster“

Typical shape of a classical CNN




Spatial resolution is decreased e.g. via max-pooling while more abstract image features are detected in deeper layers.

Building a very simple CNN with keras



```
model <- keras_model_sequential()
model %>%
  layer_conv_2d(filters=, Fill the gaps!
                kernel_size = c(5,5),
                padding = 'same',
  ...            input_shape = ...,
                activation = 'linear') %>%
  # take the max over all values in the activation map
  layer_max_pooling_2d(pool_size = ...) %>%
  layer_flatten() %>%
  layer_dense(units = 2,activation = 'softmax')
```

Exercise: Artstyle Lover



Open NB in: https://github.com/tensorchiefs/dl_course_2020/blob/master/notebooks/05_cnn_edge_lover.ipynb

Summary

- NNs work best when respecting the underlying structure of the data.
 - Use fully connected NN for tabular data
 - Use convolutional NN for data with local order such as images
- CNNs exploit the local structure of images by local connections and shared weight (same kernel is applied at each position of the image).

